

XCS234 Assignment 1

Due Sunday, August 31 at 11:59pm PT.

Guidelines

1. If you have a question about this homework, we encourage you to post your question on our Slack channel, at <http://xcs234-scpd.slack.com/>
2. Familiarize yourself with the collaboration and honor code policy before starting work.
3. For the coding problems, you must use the packages specified in the provided environment description. Since the autograder uses this environment, we will not be able to grade any submissions which import unexpected libraries.

Submission Instructions

Written Submission: Some questions in this assignment require a written response. For these questions, you should submit a PDF with your solutions online in the online student portal. As long as the PDF is legible and organized, the course staff has no preference between a handwritten and a typeset L^AT_EX submission. If you wish to typeset your submission and are new to L^AT_EX, you can get started with the following:

- Type responses only in `submission.tex`.
- Submit the compiled PDF, **not** `submission.tex`.
- Use the commented instructions within the `Makefile` and `README.md` to get started.

Coding Submission: Some questions in this assignment require a coding response. For these questions, you should submit only the `src/submission.py` file in the online student portal. For further details, see Writing Code and Running the Autograder below.

Honor code

We strongly encourage students to form study groups. Students may discuss and work on homework problems in groups. However, each student must write down the solutions independently, and without referring to written notes from the joint session. In other words, each student must understand the solution well enough in order to reconstruct it by him/herself. In addition, each student should write on the problem set the set of people with whom s/he collaborated. Further, because we occasionally reuse problem set questions from previous years, we expect students not to copy, refer to, or look at the solutions in preparing their answers. It is an honor code violation to intentionally refer to a previous year's solutions. For SCPD classes, it is also important that students avoid opening pull requests containing their solution code on the shared assignment repositories. More information regarding the Stanford honor code can be found at <https://communitystandards.stanford.edu/policies-and-guidance/honor-code>.

Writing Code and Running the Autograder

All your code should be entered into `src/submission.py`. When editing these files, please only make changes between the lines containing `### START_CODE_HERE ###` and `### END_CODE_HERE ###`. Do NOT make changes to files other than `src/submission.py`.

The unit tests in `src/grader.py` (the autograder) will be used to verify a correct submission. Run the autograder locally using the following terminal command within the `src/` subdirectory:

```
$ python grader.py
```

There are two types of unit tests used by the autograder:

- **basic:** These tests are provided to make sure that your inputs and outputs are on the right track, and can be run locally.

- **hidden:** These unit tests are NOT visible locally. These hidden tests will be run when you submit your code to the Gradescope autograder via the online student portal, and will provide feedback on how many points you have earned. These tests will evaluate elements of the assignment, and run your code with more complex inputs and corner cases. Just because your code passed the basic local tests does not necessarily mean that they will pass all of the hidden tests.

For debugging purposes, you can run a single unit test locally. For example, you can run the test case `3a-0-basic` using the following terminal command within the `src/` subdirectory:

```
$ python grader.py 3a-0-basic
```

Before beginning this course, please walk through the [Anaconda Setup for XCS Courses](#) to familiarize yourself with the coding environment. Use the env defined in `src/environment.yml` to run your code. This is the same environment used by the online autograder.

Test Cases

The autograder is a thin wrapper over the python `unittest` framework. It can be run either locally (on your computer) or remotely (on SCPD servers). The following description demonstrates what test results will look like for both local and remote execution. For the sake of example, we will consider two generic tests: `1a-0-basic` and `1a-1-hidden`.

Local Execution - Basic Tests

When a basic test like `1a-0-basic` passes locally, the autograder will indicate success:

```
----- START 1a-0-basic: Basic test case.
----- END 1a-0-basic [took 0:00:00.000062 (max allowed 1 seconds), 2/2 points]
```

When a basic test like `1a-0-basic` fails locally, the error is printed to the terminal, along with a stack trace indicating where the error occurred:

```
----- START 1a-0-basic: Basic test case.
<class 'AssertionError'>
{'a': 2, 'b': 1} != None ← This error caused the test to fail.
File "/Users/grinch/Local_Documents/Software/anaconda3/envs/XCS221/lib/python3.6/unittest/case.py", line 59, in testPartExecutor
    yield
File "/Users/grinch/Local_Documents/Software/anaconda3/envs/XCS221/lib/python3.6/unittest/case.py", line 605, in run
    testMethod()
File "/Users/grinch/Local_Documents/SCPD/XCS221/A1/src/graderUtil.py", line 54, in wrapper
    result = func(*args, **kwargs)
File "/Users/grinch/Local_Documents/SCPD/XCS221/A1/src/graderUtil.py", line 83, in wrapper
    result = func(*args, **kwargs)
File "/Users/grinch/Local_Documents/SCPD/XCS221/A1/src/grader.py", line 23, in test_0
    submission.extractWordFeatures("a b a") ← In this case, start your debugging
                                           in line 23 of grader.py.
File "/Users/grinch/Local_Documents/Software/anaconda3/envs/XCS221/lib/python3.6/unittest/case.py", line 829, in assertEquals
    assertion_func(first, second, msg=msg)
File "/Users/grinch/Local_Documents/Software/anaconda3/envs/XCS221/lib/python3.6/unittest/case.py", line 822, in _baseAssertEqual
    raise self.failureException(msg)
----- END 1a-0-basic [took 0:00:00.003809 (max allowed 1 seconds), 0/2 points]
```

Remote Execution

Basic and hidden tests are treated the same by the remote autograder, however the output of hidden tests will only appear once you upload your code to GradeScope. Here are screenshots of failed basic and hidden tests. Notice that the same information (error and stack trace) is provided as the in local autograder, now for both basic and hidden tests.

1a-0-basic) Basic test case. (0.0/2.0)

```
<class 'AssertionError': {'a': 2, 'b': 1} != None
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 59, in testPartExecutor
    yield
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 605, in run
    testMethod()
File "/autograder/source/graderUtil.py", line 54, in wrapper
    result = func(*args, **kwargs)
File "/autograder/source/graderUtil.py", line 83, in wrapper
    result = func(*args, **kwargs)
File "/autograder/source/grader.py", line 23, in test_0
    submission.extractWordFeatures("a b a"))
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 829, in assertEqual
    assertion_func(first, second, msg=msg)
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 822, in _baseAssertEqual
    raise self.failureException(msg)
```

Just like in the local autograder, this error caused the test to fail.

Just like in the local autograder, start your debugging in line 23 of grader.py.

1a-1-hidden) Test multiple instances of the same word in a sentence. (0.0/3.0)

```
<class 'AssertionError': {'a': 23, 'ab': 22, 'aa': 24, 'c': 16, 'b': 15} != None
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 59, in testPartExecutor
    yield
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 605, in run
    testMethod()
File "/autograder/source/graderUtil.py", line 54, in wrapper
    result = func(*args, **kwargs)
File "/autograder/source/graderUtil.py", line 83, in wrapper
    result = func(*args, **kwargs)
File "/autograder/source/grader.py", line 31, in test_1
    self.compare_with_solution_or_wait(submission, 'extractWordFeatures', lambda f: f(sentence))
File "/autograder/source/graderUtil.py", line 183, in compare_with_solution_or_wait
    self.assertEqual(ans1, ans2)
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 829, in assertEqual
    assertion_func(first, second, msg=msg)
File "/autograder/source/miniconda/envs/XCS221/lib/python3.6/unittest/case.py", line 822, in _baseAssertEqual
    raise self.failureException(msg)
```

This error caused the test to fail.

Start your debugging in line 31 of grader.py.

Finally, here is what it looks like when basic and hidden tests pass in the remote autograder.

1a-0-basic) Basic test case. (2.0/2.0)

1a-1-hidden) Test multiple instances of the same word in a sentence. (3.0/3.0)

0 Introduction

In this assignment, you'll delve into the core principles of Reinforcement Learning by exploring Markov Decision Processes (MDPs) and applying policy and value iteration techniques. Through a combination of coding tasks and written analyses, you'll gain hands-on experience with implementing MDPs and evaluating policies. You'll start with a coding challenge that involves solving a classic MDP problem, which will help you understand the iterative process of improving policies. You'll also examine how different decision-making horizons can affect policy outcomes, and reflect on the implications of designing reward functions, particularly in scenarios where unintended behaviors may emerge. This assignment will solidify your understanding of these foundational RL concepts and prepare you for more advanced topics.

Deliverables:

- (1) `submission.py`
- (2) Written solutions to the questions in this document

1 RiverSwim MDP

Now you will implement value iteration and policy iteration for the RiverSwim environment (see picture below¹) of (Strehl & Littman, 2008).




Figure 1: The RiverSwim MDP where dashed and solid arrows represent transitions for the LEFT and RIGHT actions, respectively. The assignment uses a modified, customizable version of what is shown above where there are three different strengths (WEAK, MEDIUM, or STRONG) of the current (transition probabilities for being pushed back or successfully swimming RIGHT).

Note: Run `python submission.py` for RiverSwim with a discount factor $\gamma = 0.99$ and a WEAK current, You can use this to verify your implementation. For grading purposes, we shall test your implementation against base and hidden test cases in `grader.py`.

(a) [4 points (Coding)]

Read through `submission.py` and implement `bellman_backup`. Return the value associated with a single Bellman backup performed for an input state-action pair.

Important: please use the version presented in the class and not any variants of it!

(b) [4 points (Coding)]

Implement `policy_evaluation` in `submission.py`. Return the optimal value function.

(c) [5 points (Coding)]

Implement `policy_improvement` in `submission.py`. Return the optimal policy.

(d) [6 points (Coding)]

Implement `policy_iteration` in `submission.py`. Return the optimal value function and the optimal policy.

(e) [6 points (Coding)]

Implement `value_iteration` in `submission.py`. Return the optimal value function and the optimal policy.

(f) [3 points (Written)]

Run both methods on RiverSwim with a WEAK current strength and find the largest discount factor (**only** up to two decimal places) such that an optimal agent starting in the initial far-left state (state s_1 in Figure 1) **does not** swim up the river (that is, does not go RIGHT). Using the value you find, interpret why this behavior makes sense. Now repeat this for RiverSwim with MEDIUM and STRONG currents, respectively. Describe and explain the changes in optimal values and discount factors you obtain both quantitatively and qualitatively.

The table above contains the values of discount factor γ (to 2 decimal places) for which the optimal agent transitioned from selecting "swim right" at the initial state s_1 ($s = 0$), to when it selected "swim left" instead. These γ values of transition were calculated for both policy iteration and value iteration dynamic programming algorithms across current strength settings WEAK, MEDIUM and STRONG of the Markov process RiverSwim.

From the table, both policy and value iteration methods have identical transition discount factors for the three different current strengths, as well as identical action-value function returns that the optimal agent compared when selecting "swim left" over "swim right" at the leftmost initial state s_1 . In RiverSwim, stronger current decreases the

¹Figure copied from (Osband & Van Roy, 2013).

Table 1: Discount factor γ and corresponding action value functions $q_*(s, a)$ for transitions to "swim left" action selection ($a = 0$) at initial state s_1 ($s = 0$)

Current strength	transition γ		$(q_*(s = 0, a = 0), q_*(s = 0, a = 1))$	
	policy iteration	value iteration	policy iteration	value iteration
WEAK	0.65	0.65	(0.0143, 0.0137)	(0.0143, 0.0137)
MEDIUM	0.76	0.76	(0.0208, 0.0203)	(0.0208, 0.0203)
STRONG	0.92	0.92	(0.0599, 0.0582)	(0.0599, 0.0582)

probability of successful "swim right" attempts in the (stochastic) dynamical state transition function T . Transition discount factors (γ) monotonically increased from WEAK to STRONG currents as the optimal agent required greater weight from the reward at the rightmost state to compensate for reduced "swim right" success probabilities in the transition function in calculations of expected future return.

2 Effect of Effective Horizon

Consider an agent managing inventory for a store, which is represented as an MDP. The stock level s refers to the number of items currently in stock (between 0 and 10, inclusive). At any time, the agent has two actions: sell (decrease stock by one, if possible) or buy (increase stock by one, if possible).

- If $s > 0$ and the agent sells, it receives +1 reward for the sale and the stock level transitions to $s - 1$. If $s = 0$ nothing happens.
- If $s < 9$ and the agent buys, it receives no reward and the stock level transitions to $s + 1$.
- The owner of the store likes to see a fully stocked inventory at the end of the day, so the agent is rewarded with +100 if the stock level ever reaches the maximum level $s = 10$.
- $s = 10$ is also a terminal state and the problem ends if it is reached.

The reward function, denoted as $r(s, a, s')$, can be summarized concisely as follows:

- $r(s, \text{sell}, s - 1) = 1$ for $s > 0$ and $r(0, \text{sell}, 0) = 0$
- $r(s, \text{buy}, s + 1) = 0$ for $s < 9$ and $r(9, \text{buy}, 10) = 100$. The last condition indicates that transitioning from $s = 9$ to $s = 10$ (fully stocked) yields +100 reward.

The stock level is assumed to always start at $s = 3$ at the beginning of the day. We will consider how the agent's optimal policy changes as we adjust the finite horizon H of the problem. Recall that the horizon H refers to a limit on the number of time steps the agent can interact with the MDP before the episode terminates, regardless of whether it has reached a terminal state. We will explore properties of the optimal policy (the policy that achieves highest episode reward) as the horizon H changes.

Consider, for example, $H = 4$. The agent can sell for three steps, transitioning from $s = 3$ to $s = 2$ to $s = 1$ to $s = 0$ receiving rewards +1, +1, and +1 for each sell action. At the fourth step, the inventory is empty so it can sell or buy, receiving no reward regardless. Then the problem terminates since time has expired.

(a) [2 points (Written)]

Starting from the initial state $s = 3$, it possible to a choose a value of H that results in the optimal policy taking both buy and sell steps during its execution? Explain why or why not.

An optimal policy is one that maximizes expected cumulative future return from the current time step to the episode horizon (Sutton & Barto, 2nd ed. p54). While the expected cumulative future return of the value function is unique in Markov Decision Processes, there may be multiple distinct policies capable of producing it; that is, an optimal policy that results from policy (or value) iteration is not necessarily unique. In the table below, two different sequences are capable of producing the same expected return with the same horizon setting $H = 4$.

For the Inventory MDP in this problem, an optimal policy can change dramatically depending on the number of steps that remain in the episode during execution before the horizon is reached. Notice in the above table that expected return jumps from 4, for $H = 6$, to 100 for $H = 7$. Also, for the same two horizon settings, execution sequences transitioned from having both *buy* and *sell* actions when $H = 6$, to *buy* only action when $H = 7$.

Inspection of the table also reveals that execution sequences under the optimal policy contain both *buy* and *sell* actions when $4 \leq H \leq 6$ and when $H \geq 8$.

(b) [2 points (Written)]

For what values of H does the optimal policy reach a fully stocked inventory, starting from the initial state $s = 3$? I.e. Give a range for H . **Note 1:** we consider the inventory fully stocked if a buy action is chosen in state $s = 9$, causing a transition to $s = 10$. This includes the last time step in the horizon. **Note 2:** By executing only buy actions, the agent can reach $s = 10$ from $s = 3$ in $H = 7$ steps.

Setting the horizon $H \geq 7$ will result in an execution sequence under the optimal policy that terminates with a fully stocked inventory. See table in 2a. Notice that a fully stocked inventory occurs regardless of whether the agent performs action selection in the terminal state s_{10} or only enters into it without performing actions in that state.

(c) [2 points (Written)]

Table 2: Table listing the expected future cumulative return without discounting ($\gamma = 1$) and the execution sequence under optimal policy across horizon values H

Horizon H	Expected Return	Execution sequence under optimal policy, $\gamma = 1$ (action, state transition, resultant reward)
4	3	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0)$
4	3	$(buy, 3 \rightarrow 4, 0), (sell, 4 \rightarrow 3, +1),$ $(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1)$
5	4	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0),$ $(sell, 0 \rightarrow 1, +1)$
6	4	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0)$
7	100	$(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, +100)$
8	100	$(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, +100), (sell, 10 \rightarrow 10, 0)$
9	101	$(sell, 3 \rightarrow 2, +1), (buy, 2 \rightarrow 3, 0),$ $(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, +100)$
10	101	$\langle$ same as above for sequence entries 1 through 8 $\rangle,$ $(buy, 9 \rightarrow 10, +100), (sell, 10 \rightarrow 10, 0)$
11	102	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(buy, 1 \rightarrow 2, 0), (buy, 2 \rightarrow 3, 0),$ $(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, 100)$
12	102	$\langle$ same as above for sequence entries 1 through 10 $\rangle,$ $(buy, 9 \rightarrow 10, 100), (sell, 10 \rightarrow 10, 0)$

Now consider the infinite-horizon discounted setting. That is, there is no time limit – the problem can only terminate when a terminal state is reached. Suppose $\gamma = 0$. What action does the optimal policy take when $s = 3$? What action does the optimal policy take when $s = 9$?

In an infinite-horizon scenario with discount factor setting $\gamma = 0$:

- When starting in state $s = 3$, the sequence of actions that the agent takes under the optimal policy is $(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1), (sell, 1 \rightarrow 0, +1), (sell, 0 \rightarrow 0, 0), (sell, 0 \rightarrow 0, 0), \dots$ ad infinitum. The sequence cannot terminate as it never exits state transition $0 \rightarrow 0$ loop caused by perpetual greedy action selection—the consequence of a myopic agent by setting $\gamma = 0$.
- For the scenario starting at $s = 9$, the execution sequence under the optimal policy migrates to the edge of the state space, in this case $s = 10$ –the terminal state–due to the inability of the agent to include non-local possibilities in calculations of expected future return. The complete execution sequence in this case is: $(buy, 9 \rightarrow 10, +100), (sell, 10 \rightarrow 10, 0)$.

(d) [2 points (Written)]

In the infinite-horizon discounted setting, is it possible to choose a fixed value of $\gamma \in [0, 1)$ such that the optimal policy starting from $s = 3$ never fully stocks the inventory? You do not need to propose a specific value, but simply explain your reasoning either way.

In the infinite-horizon scenario with reward discounting, it is possible to choose a fixed value $\gamma \in [0, 1)$ such that the optimal policy never fully stocks the inventory. Setting γ to a value for which the expected cumulative discounted return for selling surpasses that for buying at state $s = 3$ would result in the inequality

$$100 \cdot \gamma^6 < 1 + \gamma \cdot 1 + \gamma^2 \cdot 1$$

with the quantity on the left being the expected discounted future return if the agent acted to *buy* at state $s = 3$, and the quantity on the right being the expected discounted future return if the agent chose to *sell* at that state.

The numerical estimate of the real positive root for the above polynomial is $\gamma < 0.5106$. Testing about this transition γ value under value iteration with tolerance $1e^{-3}$, we find that the action selected at state $= 3$ with discount factor set to $\gamma = 0.52$ is *buy*—from which the sequence would end with the inventory fully stocked. Setting $\gamma = 0.51$, just under the theoretical value, results in selected action *sell* at state $s = 3$ whereby the agent's policy would never fully stock the inventory at the end of the day.

3 Reward Hacking

Question 2 illustrates how the particular horizon and discount factor may lead to very different policies, even with the same reward and dynamics model. This may lead to unintentional reward hacking, where the resulting policy does not match a human stakeholder's intended outcome. This problem asks you to think about an example where reward hacking may occur, introduced by Pan, Bhatia and Steinhardt². Consider designing RL for autonomous cars where the goal is to have decision policies that minimize the mean commute for all drivers (those driven by humans and those driven by AI). This reward might be tricky to specify (it depends on the destination of each car, etc) but a simpler reward (called the reward "proxy") is to maximize the mean velocity of all cars. Now consider a scenario where there is a single AI car (the red car in the figure) and many cars driven by humans (the grey car).

In this setting, under this simpler "proxy" reward, the optimal policy for the red (AI) car is to park and not merge onto the highway.³




Figure 2: Pan, Bhatia, Steinhardt ICLR 2022; <https://openreview.net/pdf?id=JYtwGwIL7ye>

(a) [2 points (Written)]

Explain why the optimal policy for the AI car is not to merge onto the highway.

To answer this question, we'll consider a scenario in which three human-operated cars (labeled by H) are traveling at the same speed along a roadway about to intercept an on-ramp on which is a single autonomous vehicle (labelled by A). We will consider a constant distance d spanning the mergepoint of the onramp onto the main roadway to another endpoint, say 1/4 mile farther down the road, and that all four vehicles will be traveling that stretch of road together without any of them turning off. The event of the autonomous vehicle merging will be labeled by M and when it is stopped on the on-ramp by S .

If the autonomous car merges onto the main roadway with velocity $v_{A,M}$ ahead of the other three human-operated cars that were traveling at $v_{H,S} > v_{A,M}$ prior to the merge event, then the human operators must reduce their speed to that of the autonomous vehicle to prevent collision. In which case $v_{H,M} = v_{A,M}$. So for the merge event mean velocity and average time of travel over distance d are respectively

$$\bar{v}_M = \frac{1}{4}(3v_{H,M} + v_{A,M}) = v_{A,M}$$

$$\bar{t}_M = \frac{d}{4}\left(\frac{3}{v_{H,M}} + \frac{1}{v_{A,M}}\right) = \frac{d}{v_{A,M}}$$

If the autonomous car is not allowed to merge so that the car's average velocity $v_{A,S} \rightarrow 0$, then correspondingly, its average time of travel over distance d becomes $t_{A,S} \rightarrow \infty$. In this case of no merge event by the autonomous vehicle, we have

²ICLR 2022 <https://openreview.net/pdf?id=JYtwGwIL7ye>

³Interestingly, it turns out that systems that use simpler function representations may reward hack less in this example than more complex representations. See Pan, Bhatia and Steinhardt's paper "The Effects of Reward Misspecification: Mapping and Mitigating Misaligned Models" for details.

$$\bar{v}_S = \lim_{v_{A,S} \rightarrow 0} \frac{1}{4}(3v_{H,S} + v_{A,S}) = \frac{3}{4}v_{H,S}$$

$$\bar{t}_S = \lim_{v_{A,S} \rightarrow 0} \frac{d}{4} \left(3 \frac{1}{v_{H,S}} + \frac{1}{v_{A,S}} \right) \approx \frac{d}{4v_{A,S}} \rightarrow \infty$$

So, if $v_{H,S} > \frac{4}{3}v_{H,M}$ then $\bar{v}_S > \bar{v}_M$, that is, the average velocity of all four cars increases even though the autonomous car remains motionless on the merge on-ramp. Notice that the average time of travel over distance d in this case $\bar{t}_S$ is much greater than the average time when the autonomous car performs the merge $\bar{t}_M$ since

$$\bar{t}_M = \frac{d}{v_{A,M}} \ll \lim_{v_{A,S} \rightarrow 0} \frac{d}{4v_{A,S}} = \bar{t}_S$$

Therefore, by using mean velocity as a proxy reward function in place of commute time in the above-described scenario (with $v_{H,S} > \frac{4}{3}v_{H,M}$), maximum mean velocity for the four cars is achieved by keeping the autonomous car stopped on the on-ramp as its merging would otherwise impede the velocity of the arterial cars and reduce the mean velocity for the entire group.

(b) [2 points (Written)]

Note this behavior is not aligned with the true reward function. Share some ideas about alternate reward functions (that are not minimizing commute) that might still be easier to optimize, but would not result in the AI car never merging. Your answer should be 2-5 sentences and can include equations: there is not a single answer and reasonable solutions will be given full credit.

Without optimizing for commute (time) which has the requirement that origin, destination, etc. of all drivers be known *a priori*, an alternative proxy reward function could be to return rewards on desirable changes in shape of the velocity distribution of all cars on the route. Returning rewards for positive shifts in median, reductions in variance and reduction of distribution weight in the near zero velocity domain could reveal policies that encourage more cars to travel at more uniform speeds. Unlike above, in which the reward function depended on a single summary statistic (e.g. the mean) that effectively promoted faster moving cars staying fast at the expense of the slower minority effectively made immobile, distribution shaping policies permit more detailed control over how potential policies can evolve by considering multiple desirable attributes in parallel.

This handout includes space for every question that requires a written response. Please feel free to use it to handwrite your solutions (legibly, please). If you choose to typeset your solutions, the —README.md— for this assignment includes instructions to regenerate this handout with your typeset $\text{\LaTeX}$ solutions.

1.f

Table 3: Discount factor γ and corresponding action value functions $q_*(s, a)$ for transitions to "swim left" action selection ($a = 0$) at initial state s_1 ($s = 0$)

Current strength	transition γ		$(q_*(s = 0, a = 0), q_*(s = 0, a = 1))$	
	policy iteration	value iteration	policy iteration	value iteration
WEAK	0.65	0.65	(0.0143, 0.0137)	(0.0143, 0.0137)
MEDIUM	0.76	0.76	(0.0208, 0.0203)	(0.0208, 0.0203)
STRONG	0.92	0.92	(0.0599, 0.0582)	(0.0599, 0.0582)

The table above contains the values of discount factor γ (to 2 decimal places) for which the optimal agent transitioned from selecting "swim right" at the initial state s_1 ($s = 0$), to when it selected "swim left" instead. These γ values of transition were calculated for both policy iteration and value iteration dynamic programming algorithms across current strength settings WEAK, MEDIUM and STRONG of the Markov process RiverSwim.

From the table, both policy and value iteration methods have identical transition discount factors for the three different current strengths, as well as identical action-value function returns that the optimal agent compared when selecting "swim left" over "swim right" at the leftmost initial state s_1 . In RiverSwim, stronger current decreases the probability of successful "swim right" attempts in the (stochastic) dynamical state transition function T . Transition discount factors (γ) monotonically increased from WEAK to STRONG currents as the optimal agent required greater weight from the reward at the rightmost state to compensate for reduced "swim right" success probabilities in the transition function in calculations of expected future return.

2.a

An optimal policy is one that maximizes expected cumulative future return from the current time step to the episode horizon (Sutton & Barto, 2nd ed. p54). While the expected cumulative future return of the value function is unique in Markov Decision Processes, there may be multiple distinct policies capable of producing it; that is, an optimal policy that results from policy (or value) iteration is not necessarily unique. In the table below, two different sequences are capable of producing the same expected return with the same horizon setting $H = 4$.

Table 4: Table listing the expected future cumulative return without discounting ($\gamma = 1$) and the execution sequence under optimal policy across horizon values H

Horizon H	Expected Return	Execution sequence under optimal policy, $\gamma = 1$ (action, state transition, resultant reward)
4	3	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0)$
4	3	$(buy, 3 \rightarrow 4, 0), (sell, 4 \rightarrow 3, +1),$ $(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1)$
5	4	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0),$ $(sell, 0 \rightarrow 1, +1)$
6	4	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0),$ $(sell, 0 \rightarrow 1, +1), (buy, 0 \rightarrow 1, 0)$
7	100	$(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, +100)$
8	100	$(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, +100), (sell, 10 \rightarrow 10, 0)$
9	101	$(sell, 3 \rightarrow 2, +1), (buy, 2 \rightarrow 3, 0),$ $(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, +100)$
10	101	$\langle$ same as above for sequence entries 1 through 8 $\rangle,$ $(buy, 9 \rightarrow 10, +100), (sell, 10 \rightarrow 10, 0)$
11	102	$(sell, 3 \rightarrow 2, +1), (sell, 2 \rightarrow 1, +1),$ $(buy, 1 \rightarrow 2, 0), (buy, 2 \rightarrow 3, 0),$ $(buy, 3 \rightarrow 4, 0), (buy, 4 \rightarrow 5, 0),$ $(buy, 5 \rightarrow 6, 0), (buy, 6 \rightarrow 7, 0),$ $(buy, 7 \rightarrow 8, 0), (buy, 8 \rightarrow 9, 0),$ $(buy, 9 \rightarrow 10, 100)$
12	102	$\langle$ same as above for sequence entries 1 through 10 $\rangle,$ $(buy, 9 \rightarrow 10, 100), (sell, 10 \rightarrow 10, 0)$

For the Inventory MDP in this problem, an optimal policy can change dramatically depending on the number of steps that remain in the episode during execution before the horizon is reached. Notice in the above table that expected return jumps from 4, for $H = 6$, to 100 for $H = 7$. Also, for the same two horizon settings, execution sequences transitioned from having both *buy* and *sell* actions when $H = 6$, to *buy* only action when $H = 7$.

Inspection of the table also reveals that execution sequences under the optimal policy contain both *buy* and *sell* actions when $4 \leq H \leq 6$ and when $H \geq 8$.

2.b

Setting the horizon $H \geq 7$ will result in an execution sequence under the optimal policy that terminates with a fully stocked inventory. See table in 2a. Notice that a fully stocked inventory occurs regardless of whether the agent performs action selection in the terminal state s_{10} or only enters into it without performing actions in that state.

2.c

In an infinite-horizon scenario with discount factor setting $\gamma = 0$:

- When starting in state $s = 3$, the sequence of actions that the agent takes under the optimal policy is $(sell, 3 \rightarrow 2, +1)$, $(sell, 2 \rightarrow 1, +1)$, $(sell, 1 \rightarrow 0, +1)$, $(sell, 0 \rightarrow 0, 0)$, $(sell, 0 \rightarrow 0, 0)$, ... ad infinitum. The sequence cannot terminate as it never exits state transition $0 \rightarrow 0$ loop caused by perpetual greedy action selection—the consequence of a myopic agent by setting $\gamma = 0$.
- For the scenario starting at $s = 9$, the execution sequence under the optimal policy migrates to the edge of the state space, in this case $s = 10$ —the terminal state—due to the inability of the agent to include non-local possibilities in calculations of expected future return. The complete execution sequence in this case is: $(buy, 9 \rightarrow 10, +100)$, $(sell, 10 \rightarrow 10, 0)$.

2.d

In the infinite-horizon scenario with reward discounting, it is possible to choose a fixed value $\gamma \in [0, 1)$ such that the optimal policy never fully stocks the inventory. Setting γ to a value for which the expected cumulative discounted return for selling surpasses that for buying at state $s = 3$ would result in the inequality

$$100 \cdot \gamma^6 < 1 + \gamma \cdot 1 + \gamma^2 \cdot 1$$

with the quantity on the left being the expected discounted future return if the agent acted to *buy* at state $s = 3$, and the quantity on the right being the expected discounted future return if the agent chose to *sell* at that state.

The numerical estimate of the real positive root for the above polynomial is $\gamma < 0.5106$. Testing about this transition γ value under value iteration with tolerance $1e^{-3}$, we find that the action selected at state $= 3$ with discount factor set to $\gamma = 0.52$ is *buy*—from which the sequence would end with the inventory fully stocked. Setting $\gamma = 0.51$, just under the theoretical value, results in selected action *sell* at state $s = 3$ whereby the agent's policy would never fully stock the inventory at the end of the day.

3.a

To answer this question, we'll consider a scenario in which three human-operated cars (labeled by H) are traveling at the same speed along a roadway about to intercept an on-ramp on which is a single autonomous vehicle (labelled by A). We will consider a constant distance d spanning the merge point of the onramp onto the main roadway to another endpoint, say 1/4 mile farther down the road, and that all four vehicles will be traveling that stretch of road together without any of them turning off. The event of the autonomous vehicle merging will be labeled by M and when it is stopped on the on-ramp by S .

If the autonomous car merges onto the main roadway with velocity $v_{A,M}$ ahead of the other three human-operated cars that were traveling at $v_{H,S} > v_{A,M}$ prior to the merge event, then the human operators must reduce their speed to that of the autonomous vehicle to prevent collision. In which case $v_{H,M} = v_{A,M}$. So for the merge event mean velocity and average time of travel over distance d are respectively

$$\bar{v}_M = \frac{1}{4}(3v_{H,M} + v_{A,M}) = v_{A,M}$$

$$\bar{t}_M = \frac{d}{4}\left(\frac{3}{v_{H,M}} + \frac{1}{v_{A,M}}\right) = \frac{d}{v_{A,M}}$$

If the autonomous car is not allowed to merge so that the car's average velocity $v_{A,S} \rightarrow 0$, then correspondingly, its average time of travel over distance d becomes $t_{A,S} \rightarrow \infty$. In this case of no merge event by the autonomous vehicle, we have

$$\bar{v}_S = \lim_{v_{A,S} \rightarrow 0} \frac{1}{4}(3v_{H,S} + v_{A,S}) = \frac{3}{4}v_{H,S}$$

$$\bar{t}_S = \lim_{v_{A,S} \rightarrow 0} \frac{d}{4}\left(3\frac{1}{v_{H,S}} + \frac{1}{v_{A,S}}\right) \approx \frac{d}{4v_{A,S}} \rightarrow \infty$$

So, if $v_{H,S} > \frac{4}{3}v_{H,M}$ then $\bar{v}_S > \bar{v}_M$, that is, the average velocity of all four cars increases even though the autonomous car remains motionless on the merge on-ramp. Notice that the average time of travel over distance d in this case $\bar{t}_S$ is much greater than the average time when the autonomous car performs the merge $\bar{t}_M$ since

$$\bar{t}_M = \frac{d}{v_{A,M}} \ll \lim_{v_{A,S} \rightarrow 0} \frac{d}{4v_{A,S}} = \bar{t}_S$$

Therefore, by using mean velocity as a proxy reward function in place of commute time in the above-described scenario (with $v_{H,S} > \frac{4}{3}v_{H,M}$), maximum mean velocity for the four cars is achieved by keeping the autonomous car stopped on the on-ramp as its merging would otherwise impede the velocity of the arterial cars and reduce the mean velocity for the entire group.

3.b

Without optimizing for commute (time) which has the requirement that origin, destination, etc. of all drivers be known *a priori*, an alternative proxy reward function could be to return rewards on desirable changes in shape of the velocity distribution of all cars on the route. Returning rewards for positive shifts in median, reductions in variance and reduction of distribution weight in the near zero velocity domain could reveal policies that encourage more cars to travel at more uniform speeds. Unlike above, in which the reward function depended on a single summary statistic (e.g. the mean) that effectively promoted faster moving cars staying fast at the expense of the slower minority effectively made immobile, distribution shaping policies permit more detailed control over how potential policies can evolve by considering multiple desirable attributes in parallel.